

Device-Anchored Irrigation Logs: An Adversarial Evaluation of Near-Zero-Cost MRV for Low-Emission Rice

Xuan Van Mai¹, Tuong Tri Nguyen^{2,3,*}

¹ Thuan Hoa High School, University of Education, Hue University, Hue, Viet Nam

² Institute of Open Education and Information Technology, Hue University, Hue, Viet Nam

³ University of Education, Hue University, 34 Le Loi St., Hue, Viet Nam

* Corresponding author: Tuong Tri Nguyen. E-mail: ntuongtri@hueuni.edu.vn

Abstract. Monitoring, reporting and verification (MRV) is the bottleneck keeping emission reductions in low-emission rice from becoming tradable carbon credits, and monitoring is about three quarters of that cost. Because every credit-eligible action under alternate wetting and drying is already an *actuation command*, a farmer-written irrigation log looks like free evidence of the highest rank under Decision 4801/QD-BNNMT. We built that layer on an FAO-56 digital twin, then attacked it; the attack is the contribution. Replaying a log under public ERA5 forcing exposes claimed depths no soil can hold, but an adversary who keeps the total volume and every individual depth and merely *shifts commands in time* inflates creditable dry phases by 33% in paired comparison. No physical tier escapes it usefully: the point model accuses 26.7% of honest farmers, and the robust tier that cuts false accusation to 5.0% in turn misses most time-shifting attacks. Over 281 audited logs at three Mekong Delta stations and two climate years, matching the log against the gateway's signed, time-stamped command record catches 116/116 forged logs and accuses none. Physics-based attestation cannot stand alone; the timestamp must come from a device, not the farmer.

Keywords: alternate wetting and drying; monitoring, reporting and verification; carbon credit; low-emission rice; digital twin.

1 Introduction

Viet Nam has committed to one million hectares of low-emission rice by 2030, and the agronomy is settled: raising the severity of intermittent drying cut methane emission by 41–71% against continuous flooding with no grain-yield penalty [1], and water regime dominates seasonal methane flux across Asian rice fields [2]. What is not settled is measurement. Under Decision 4801/QD-BNNMT, documentary evidence for a credit claim is ranked, and the record of the irrigation and drainage actions actually taken sits at the top of that ranking (spot-check Form 3.3, Annex 2: actual irrigation and drainage records above plans, schedules, field notes and signed statements) [3], and under Verra's VM0051 the same record must survive third-party verification [4]. Producing it costs money a smallholder cooperative does not have: monitoring alone is about 77% of MRV expenditure [5], and the current procedure relies on periodic photographs of a water-level tube.

The obvious repair is to notice that a modern irrigation command is already a digital event: if a gateway actuates the pump it can also record the actuation, and that record is a Tier-1 activity datum in the sense of Decision 4801 at no marginal cost — the premise of the companion work [6],

which uses an FAO-56 twin as a safety gate deciding whether a command may be transmitted over a lossy channel. We took the next step and asked whether the same twin could *audit* what a farmer claims after the fact, with no sensor deployed: a log of claimed depths is replayed through the water balance under public ERA5 forcing, and a claim the soil could not have absorbed is evidence of fabrication. A first version of this layer reported no false accusation on 60 honest logs and detection of 60/60 forged ones — correct and worthless, because the trajectories labelled “real” were generated by the same model that then audited them: the detector was grading its own homework (Supplementary S10.4).

This paper reports what happens when that circularity is removed: we rebuild the field so that it violates every assumption of the twin, and replace the naive fraud patterns with an adversary that has read the detector. Three findings follow, and they are also the contributions. First, the point-model detector is not merely optimistic, it is *unsafe*: it falsely accuses 26.7% of honest farmers, because an absolute pump-gap threshold treats a farmer who forgets one entry exactly like one who irrigated secretly; the robust rule that fixes this — an intersection over an uncertainty set Θ of 30 soil, crop-coefficient and efficiency configurations (Sec. 6.2) — pays for it in detection power, which we measure rather than hide. Second, physics-based attestation has a hole no amount of modelling closes: an adversary who preserves the total volume and every individual depth can *shift commands in time*, inflating the credit-eligible dry phases that VM0051 and Decision 4801 pay for by +33% in paired comparison ($t = 5.29$; +38%, $t = 3.46$, full uncertainty set). No physical tier escapes — the point model catches four of nine altered logs but is the very detector that accuses 26.7% of honest farmers, while the robust tier that cuts false accusation to 5.0% lets eight of nine through — and we prove this is unavoidable for any detector that is a function of the replayed trajectory (Sec. 3, Sec. 5). Third, the repair is cheap and not physical: matching the log against the gateway’s signed, time-stamped command record catches 116/116 genuinely forged logs and accuses none of 60 honest ones, at no cost beyond the gateway the companion work deploys (Sec. 7). The paper also corrects three published numbers: the admissible depth is a range [49.0, 86.0] mm, not 61 mm; the third physical filter cannot activate on dry-season delta rice; and the absolute pump-gap threshold is inadmissible (Sec. 6.1, Sec. 1).

2 System Model

2.1 Field water balance

The root-zone water content w_t [mm] evolves hourly under the FAO-56 bucket model [7]:

$$w_{t+1} = \text{clip}_{\leq FC}(w_t + P_t - RO_t - ET_{c,t} + u_t), \quad (1)$$

with crop evapotranspiration and a water-stress coefficient

$$ET_{c,t} = K_{s,t} K_c ET_{0,t}, \quad K_{s,t} = \frac{FC - w_t}{(1 - p) TAW}, \quad (2)$$

where $ET_{0,t}$ is hourly reference evapotranspiration (Penman-Monteith), P_t rainfall, RO_t surface runoff (Soil Conservation Service [SCS] curve-number, evaluated daily and shared pro rata over

the hours of a rain event), and u_t the irrigation depth applied in hour t . Drainage above field capacity FC is removed by the clipping operator, so mass conservation holds exactly and is checked as an invariant. The root-zone depletion fraction is

$$d_t = \frac{FC - w_t}{FC - WP}, \quad \text{span} = FC - WP, \quad (3)$$

with wilting point WP ; for the reference silty clay loam of the delta, $TAW = FC - WP = 70$ mm, $FC = 160$ mm, $WP = 90$ mm.

2.2 Replay of a claimed log

Let a claimed log supply the irrigation term of (1). The *replay operator* $\mathcal{R}$ maps a log to the trajectory it implies, by iterating (1) under the observed forcing $(P_t, ET_{0,t})$:

$$\mathcal{R}(L) = \{(\hat{w}_t, \hat{d}_t)\}_{t=0}^T, \quad (4)$$

with $\hat{w}_0 = w_{\text{init}}$ and

$$u_t = \sum_{(t_i, u_i) \in L} u_i \mathbb{1}[t_i = t], \quad (5)$$

so that $\hat{w}_t$ is the root-zone water content the log *claims* the field had, and $\hat{d}_t$ its depletion fraction under (3); the hat marks quantities reconstructed from the log rather than measured. Replay is deterministic and uses only public reanalysis, so two auditors replaying the same log obtain the same trajectory and the verdict is reproducible. A superscript minus denotes the state immediately *before* that hour's claim is applied, as in $\hat{w}_t^-$: the water available to receive an event, which is what a filter must inspect.

The operator is also what makes the scheme cheap: no sensor is queried, and the only inputs are the log, the public weather series and a soil profile. Its weakness, developed in Sec. 5, is that it can only check whether a trajectory is *possible*, never whether it *happened*.

3 Threat Model

A claimed irrigation log is a finite map

$$L = \{(t_i, u_i)\}_{i=1}^{|L|}, \quad u_i > 0, \quad t_i \in [0, T], \quad (6)$$

where u_i is a claimed depth in millimetres and t_i an hour index. Three independent quantities can be forged, and each defence tier observes a different subset of them:

$$\underbrace{\sum_i u_i}_{\text{total volume}}, \quad \underbrace{\{u_i\}}_{\text{individual depths}}, \quad \underbrace{\{t_i\}}_{\text{timestamps}}. \quad (7)$$

Supplementary Table S1 lists the resulting fraud classes. The adversary is *rational*: it maximises creditable dry phases subject to not being flagged, and it knows the detector, because the detector is published; this is the sense in which ε in (12) below is not a secret.

Let M be the volume actually pumped, measured at the cooperative station, and $U(L) = \sum_i u_i$ the claimed total; the tier-1 statistic is

$$G = M - U(L), \quad (8)$$

and the credit-relevant quantity counts *dry phases*. Write $A(\cdot)$ for the set of maximal intervals $[a, b]$ of a replayed trajectory over which the depletion stays at or above d_{dry} with no claimed irrigation, lasting at least 72 h — the three-day minimum of Decision 4801 [3]:

$$\begin{aligned} A(\{\hat{d}_t\}) = \{[a, b] : & \min_{t \in [a, b)} \hat{d}_t \geq d_{\text{dry}}, \\ & u_t = 0 \ \forall t \in [a, b), \\ & b - a \geq 72 \text{ h}\}. \end{aligned} \quad (9)$$

The credited count is then

$$n_{\text{dry}}(L) = |A(\mathcal{R}(L))|, \quad (10)$$

with $d_{\text{dry}} = 0.45$ throughout. Revenue is proportional to n_{dry} , not to $U(L)$ — the distinction that makes the attack of Sec. 5 profitable, invisible in a formulation that treats total volume as the target.

4 Physics-Based Filters and Their Exact Blind Spots

Replaying L under public ERA5 forcing produces $(\hat{w}_t, \hat{d}_t)$ (the companion work [6] gives the balance in full); three filters apply. V_1 , *saturation*: a claim issued while the root zone is already at field capacity cannot be stored:

$$V_1(L) = \{i : u_i > 0 \wedge \hat{w}_{t_i}^- \geq FC - \varepsilon\}. \quad (11)$$

V_4 , *storage excess*: a claim larger than the deficit available at that hour must have run off or percolated, so it was not “applied” in the sense the standard requires:

$$V_4(L) = \{i : u_i > \underbrace{(FC - \hat{w}_{t_i}^-)}_{\text{deficit}} + \varepsilon\}. \quad (12)$$

V_5 , *infiltration bound*. Independent of state, a single event of duration τ_i cannot deliver more than the soil can accept:

$$u_i \leq k_{\text{inf}} \tau_i + p_{\text{max}}, \quad (13)$$

with $k_{\text{inf}} = 10 \text{ mm h}^{-1}$ and $p_{\text{max}} = 100 \text{ mm}$ at the loose end of the published range for puddled rice (bound 110 mm at $\tau_i = 1 \text{ h}$), so that a farmer who records two adjacent events in one hour is not accused.

A fourth filter, on depletion below the wilting point, was included in the published version of this layer. It cannot activate: see Sec. 1.

Proposition 1 (Exact blindness of aggregate defences). *Let L and L' share the same multiset of depths and the same total, i.e. $\{u_i\} = \{u'_i\}$ as multisets. Then*

$$G(L) = G(L') \quad \text{and} \quad |V_5(L)| = |V_5(L')|. \quad (14)$$

Proof. Supplementary S11; both quantities are symmetric functions of the depth multiset and never depend on t_i . $\square$

Proposition 1 is the whole difficulty in one line: the two defences that do not depend on the model being right — the pump meter and the infiltration bound — are *exactly* invariant under the transformation the adversary needs. What remains is V_1 and V_4 , which do depend on order through $\hat{w}_{t_i}^-$; they turn out not to help.

5 The Time-Shifting Attack

Proposition 2 (Credit gain and evasion are the same action). *Let a claim at hour t have available deficit $c_t = FC - \hat{w}_t^-$. Between consecutive claims the bucket loses $\Delta_t = \sum_s (ETc_s + DP_s - P_s^+) \geq 0$, so delaying a claim by δ hours increases its deficit by the drainage accrued in those hours:*

$$c_{t+\delta} \geq c_t + \min(\delta \cdot \bar{\Delta}, FC - \hat{w}_t^-). \quad (15)$$

Consequently the transformation that moves claims out of a dry interval and onto its boundary simultaneously (i) increases n_{dry} , because the interval is no longer interrupted, and (ii) increases c_t at every moved claim, which by (12) makes V_4 less likely to fire. There is no trade-off for the adversary to negotiate: the gradient of profit and the gradient of undetectability point the same way.

Proof. Supplementary S11. In short: (i) is definitional, since A requires no claim inside the interval; (ii) follows because delaying a claim lets the bucket drain for δ further hours with no inflow, so the available deficit can only grow while u_i is unchanged. $\square$

This is why the attack cannot be closed by better hydrology: any detector whose verdict is a function of the replayed trajectory must integrate the same drainage that makes the delayed claim look legitimate. The empirical attack is a hill-climb over integer hour offsets in $[-36, 36]$ accepting a move only if the total is preserved bit-for-bit, every depth stays under (13), and n_{dry} strictly increases. The attack does not steal water — what is stolen is *history* (Supplementary S19).

6 When May a Farmer Be Accused?

6.1 The admissible depth is a range, not a number

The published version of this layer stated a single bound, $u \leq 61$ mm, as the largest claim that never raises an alarm — an artefact of one soil profile and one initial condition. With V_4 firing when $u_i > c_{t_i} + \varepsilon$ and $c_t = d_t \cdot \text{span}$, the bound at a field with initial depletion w_0 irrigating no later than d^* is

$$u_{\max} = \max(d^*, w_0) \cdot \text{span} + \varepsilon. \quad (16)$$

Two corrections follow, both verified by scanning the replay operator rather than by algebra alone (derivation in Supplementary S14): (16) uses max, not a difference — an initially dry field admits *larger* claims than a wet one, and the previously published expression corresponds to no physical

state — and it bounds a *single event*, not the season, so an adversary issuing many sub-threshold claims can inflate the total without bound. Supplementary Table S2 gives the range that must be published instead, [49.0, 86.0] mm across Θ .

6.2 Robust accusation under model uncertainty

The auditor knows neither the farmer’s soil, nor the growth stage, nor the delivery efficiency. Let Θ be the uncertainty set and $V_\theta(L)$ the violation set under configuration θ ; the published detector evaluated a single default θ_0 and treated every violation as guilt. We replace it by an intersection:

$$V_{\text{rob}}(L) = \bigcap_{\theta \in \Theta} V_\theta(L), \quad (17)$$

$$\Theta = \mathcal{S} \times \mathcal{K} \times \mathcal{E} \times \mathcal{W},$$

with $|\mathcal{S}| = 5$ soil textures, $|\mathcal{K}| = 2$ crop-coefficient regimes, $|\mathcal{E}| = 3$ irrigation efficiencies and $|\mathcal{W}| = 5$ initial states, so $|\Theta| = 150$ and each log costs 150 replays. The semantics are deliberately one-sided:

$$a(L) = \begin{cases} \text{accuse,} & V_{\text{rob}}(L) \neq \emptyset, \\ \text{query,} & \text{otherwise.} \end{cases} \quad (18)$$

Proposition 3 (No false accusation within Θ). *If the true field configuration θ^* belongs to Θ and L is compliant under θ^* , then $V_{\text{rob}}(L) = \emptyset$ and the farmer is not accused.*

Proof. $V_{\text{rob}}(L) \subseteq V_{\theta^*}(L) = \emptyset$. □

Proposition 3 is what makes the layer deployable: model error can reduce detection power but cannot manufacture an accusation, at a price that must be stated. Two implementation details — the treatment of non-firing worlds and the span of $\mathcal{W}$ — silently break the proposition if got wrong (Supplementary S7).

6.3 A proportional pump-gap threshold

The absolute rule $G > \tau_0$ with $\tau_0 = 5$ mm assumes the log is a faithful copy of what was pumped. It is not: forgetting, rounding to the nearest 5 mm and day-late entries all contribute to G without wrongdoing, so the threshold must scale with the number of recorded events:

$$\tau(n) = \delta_{\text{rel}} M + z\sqrt{n} \sigma_{\text{event}}, \quad (19)$$

$$\sigma_{\text{event}}^2 = p(1-p)\bar{u}^2 + \frac{r^2}{12} + q(1-q)(s\bar{u})^2,$$

with p the forgetting rate, r the rounding step, q the depth-error rate and s its amplitude, $\bar{u} = 30$ mm a typical event depth, $\delta_{\text{rel}} = 2\%$ and $z = 3$; the absolute rule is the degenerate case $n \rightarrow 0$, $\sigma_{\text{event}} \rightarrow 0$.

Remark 1 (A filter that cannot fire). *The published layer’s third physical filter, flagging depletion below the wilting point, was doubly inoperative: its threshold (1.02) lay outside the clipped range of the reported indicator (0 firings in 281 logs), and once repaired it still cannot activate — across 150 configurations with an empty log the unclipped depletion peaks at 0.9987, and a dry-season delta rice field does not reach the*

wilting point (Supplementary Table S3, derivation in S13). The layer has two operative physical filters and one inert safety device.

7 Device-Anchored Attestation

7.1 The gateway record

Propositions 1–2 show that volume and depth are insufficient and that order is exactly what the adversary manipulates, so the defence must observe timestamps — and the only party that can vouch for a timestamp is the device that acted on it. Let the gateway hold a signed command record

$$\mathcal{G} = \{(t_j, u_j, \sigma_j)\}, \quad (20)$$

where σ_j is a signature over (t_j, u_j) produced by the controller at actuation time. The anchor test compares the claimed log against it:

$$\begin{aligned} \Delta_{\text{time}} &= \{t_i\} \Delta \{t_j\}, \\ \Delta_{\text{depth}} &= \max_t |u(t) - u_{\mathcal{G}}(t)|, \end{aligned} \quad (21)$$

where Δ is the symmetric difference and Δ_{depth} is taken over hours common to both records. The verdict is then

$$A(L, \mathcal{G}) = \begin{cases} \text{forge}, & \Delta_{\text{time}} \neq \emptyset \vee \Delta_{\text{depth}} > \epsilon, \\ \text{match}, & \text{otherwise}, \end{cases} \quad (22)$$

with $\epsilon = 10^{-6}$ mm, i.e. exact agreement up to floating-point noise. Naming the two discrepancies separately matters: time-shifting alters Δ_{time} only, depth inflation alters Δ_{depth} only, and omission and addition alter both, so a single comparison catches every fraud class.

Proposition 4 (Completeness of the anchor). *If a claimed log L differs from the gateway record $\mathcal{G}$ in the multiset of hours or in any depth at a shared hour, then $A(L, \mathcal{G})$ returns forge. If L and $\mathcal{G}$ agree on both, then L is a faithful transcript and no forgery of the three quantities in (7) has occurred.*

Proof. Supplementary S12: the first clause is the definition of (22), and the second follows because all three forgeable quantities are functions of the hour set and the per-hour depths. $\square$

7.2 Why this does not re-introduce false accusation

Proposition 4 is strong enough to look dangerous: any discrepancy is treated as forgery, including the forgetting and rounding that Sec. 6 showed to be ubiquitous. The resolution is that the anchor compares a transcript against the *source*, not a log against a *model*: an honest farmer who forgets an entry is simply completed from $\mathcal{G}$. False accusation then requires the gateway to disagree with itself — a device-integrity question — which is why Table 1 holds at 0/60 while the behavioural tiers hold at 26.7% and 5.0% (Supplementary S18).

7.3 Cost and the scope of the claim

The anchor requires the gateway of the companion work [6], which actuates the pump and already writes a signed, time-stamped command record as part of its safety function; the marginal cost of attestation is storage and transmission of a record that exists regardless (Supplementary Table S5). The scope is nonetheless narrower than the premise we started from: a hand-written log, with no device in the loop, is *not* admissible evidence on its own, since Sec. 9.2 measures a rational adversary inflating credits by +35% while escaping every physics- and volume-based test (6/6 altered logs). Near-zero-cost MRV is achievable only where a controller stands between the farmer and the pump; where irrigation is manual this layer does not apply.

8 Experimental Design

The design compares four defence layers against nine log types on a field that deliberately violates every assumption of the twin: unlike the prior evaluation, whose 0/60 false accusations measured only self-consistency (Supplementary S10.4), the field here is generated by a separate process that differs from the twin along seven axes, each with a published source (every value in Supplementary Table S6; axis-by-axis list and per-event recording semantics in S20.1). Weather is hourly ERA5 (Open-Meteo archive) for Can Tho, Soc Trang and Ca Mau over 2024 (water-scarce) and 2025 (water-surplus), window 1 March to 29 May (2 160 h), the dry season in which AWD is practised. Layer 1 is the published point-model detector (default profile, absolute threshold $\tau_0 = 5$ mm); Layer 2 the robust detector of (17) with the proportional threshold of (19), on both a reduced Θ (24 replays per log, 10 seeds) and the full one (150 replays, cross-check); Layer 3 the device anchor of (22); Layer 4 physics alone, for reference. Of the nine log types, three are naive forgeries and four adversarial, built by an optimiser rather than a fixed recipe (full list in S16); all attack statistics are conditioned on a changed flag, because an adversary that finds no improving move returns the original log, and comparisons use the baseline of the variant itself (two harness defects, S10.2). For each layer we report the escape rate (forged logs neither physical tier flags) and the false-accusation rate on honest logs; for credit impact, n_{dry} of (10) against the honest field baseline, paired. Volume appears only as a diagnostic (S10.3).

9 Results

9.1 The published detector is unsafe, not merely optimistic

Table 1 gives the central negative result: the point-model detector accuses nobody on paper-condition logs, reproducing the prior 0/60, but accuses 26.7% of honest farmers on field-condition logs. The cause is (8) with an absolute threshold — forgetting, rounding or a day-late entry all produce $G > 5$ mm without wrongdoing; the measured gap on honest field-condition logs reached 65 mm. The proportional threshold of (19) cuts this to 5.0%, the device anchor to 0/60.

Table 1. False accusation on honest logs, the property the scheme is sold on. Paper conditions reproduce the prior 0/60; field conditions expose the absolute gap threshold.

Defence layer	Honest, paper conditions	Honest, field conditions
v1 tier 0 (physics, point model)	0/60	2/60
v1 tier 1 (absolute $G > 5$ mm)	0/60	16/60
v2 tier 0 ($\forall$ over Θ)	0/60	0/60
v2 tier 1 ($\tau(n)$ scales with events)	0/60	3/60
Device-anchored command log	0/60	0/60

9.2 Physics cannot see the attack that pays

Table 2 is the second negative result, and the one that matters economically. Every inflation figure is a *paired* comparison against the honest log of the same trajectory, restricted to attacks that actually changed the record; the unpaired alternative overstates the gain roughly twofold (S15). Time-shifting raises creditable dry phases from 4.78 to 6.33 per season, +33% over nine altered logs ($t = 5.29$); under the full uncertainty set, +38% over three ($t = 3.46$, $p = 0.074$ — directionally consistent but not significant at 0.05 with this sample); the rational adversary reaches +35% and +40% ($p = 0.007$ and $p = 0.30$). Consolidation gains almost nothing once paired (+1%, +2%), and the two crude forgeries *reduce* creditable phases (−1%, −17%): only the volume-preserving permutation of timestamps is profitable, as Proposition 2 predicts. The escape rates in Table 3 are the heart of the result: time-shifting escapes the robust physics layer in 8 of 9 altered logs, the rational adversary in 6 of 6. The point model catches more (4 of 9), but it is precisely the detector that accuses 26.7% of honest farmers; the robust tier cuts that to 5.0% and hands most of the detection back (1 of 9) — *no physical configuration escapes that trade-off*, because both tiers are functions of the replayed trajectory. The infiltration filter V_5 fires 648 times, but exclusively on the crude depth-inflation forgery: useful, and orthogonal to the threat that pays.

9.3 The anchor closes every class, and the cost of robustness

Table 3 also gives the constructive half: of 281 audited logs, 116 were genuinely altered, and the device anchor caught 116/116 while accusing none of 60 honest field-condition logs (37/37 and 0/18 under the full Θ). Naive forgeries are caught by every layer; the distinction lies in the adversarial rows, where only the anchor holds. Robustness is not free: on the nine omission-fraud logs actually altered the point model catches 9/9, the robust rule 4/9 (4/4 vs 2/4 full Θ) — a loss of detection power, not a Pareto improvement, though the correct side of the trade. Every conclusion holds under both the reduced Θ (24 replays, 281 logs, 10 seeds) and the full one (150 replays, 85 logs, 3 seeds): false accusation 0/60 and 0/18, anchor completeness 116/116 and 37/37, time-shift escape 8/9 and 3/3, paired inflation +33% and +38% (Supplementary S17).

Table 2. Creditable dry phases of at least 72 h — the quantity VM0051 and Decision 4801 pay for — against the honest field baseline (paired). The last column counts attacks that actually altered the log. The table carries two baselines, deliberately different: the inflation column uses the paired baselines against the honest log of the same station, year, seed and soil (4.87 for the adversarial variants), while the “honest” row reports the mean 3.23 phases over all 60 honest logs. See Supplementary S15.

Log type	Altered logs	Paired dry phases (before → after)	Inflation	t
Honest farmer (field conditions)	—	3.23 (reference)	reference	—
Added irrigation events	23/23	4.87 → 4.83	-1%	-0.57
Infeasible claimed depth	23/23	4.87 → 4.04	-17%	-4.03
Omission of real commands	9/23	5.44 → 4.78	-12%	-2.83
Claim consolidation, 24 h	23/23	4.87 → 4.91	+1%	0.37
Claim consolidation, 96 h	23/23	4.87 → 4.96	+2%	0.62
Time shifting (rational)	9/23	4.78 → 6.33	+33%	5.29
Rational adversary	6/23	4.33 → 5.83	+35%	4.39

Table 3. Main result. Rates are computed over the logs an attack *actually altered*; column 2 reports altered over generated, since the adversary returns the original log when it finds no improving move. Escape rate is the fraction neither physical tier flags.

Fraud variant	Altered logs	Escapes v1	Escapes v2 physics	Caught by anchor
Honest farmer (field conditions)	0/60 (honest)	27% accused	5% accused	0% accused
Added irrigation events	23/23	0%	0%	23/23
Infeasible claimed depth	23/23	0%	0%	23/23
Omission of real commands	9/23	0%	56%	9/9
Claim consolidation, 24 h	23/23	39%	91%	23/23
Claim consolidation, 96 h	23/23	39%	91%	23/23
Time shifting (rational)	9/23	56%	89%	9/9
Rational adversary	6/23	100%	100%	6/6
All genuinely forged	116	—	—	116/116

Table 4. Confirmable credit-eligible dry phases under each MRV route, over the 150 phases present in the 60 simulated fields; a phase counts only if at least one observation falls inside it. The final two rows tie here and differ in cost and in what they prove about a doctored log.

Verification route	Confirmable dry phases ($n = 150$)	Observations per season	Devices per plot
Photo + self-report (current practice) [3]	0/150 (0%)	—	0 (PVC tube)
Sentinel-1A+1B (6-day revisit) [8]	124/150 (83%)	15	0
Sentinel-1A alone (12-day revisit) [9]	90/150 (60%)	8	0
Continuous water-level tube (IoT) [10]	150/150 (100%)	2,153	1 per plot
Signed command record (this work) [6]	150/150 (100%)	2,153	0 (controller)

10 Discussion

10.1 Relation to prior work

AWD monitoring has been approached from three directions — instrumented water-level sensors [10], remote sensing [11–13], and programme-level MRV [14–16] — while agronomic meta-analysis [17] and adoption studies [18, 19] explain why instrumented monitoring has not spread. None evaluates the *reliability of the action record itself*, which is the question asked here (full survey in Supplementary S8). Satellite radar, the obvious low-cost alternative, we measured rather than argued about: Table 4 scores every route by how many of the 150 credit-eligible dry phases it could confirm, a phase counting only if at least one observation falls inside it [8–10]. Two cautions: the column measures the ability to *earn* a credit, not to detect fraud — tube and command record tie at 100% but differ in cost and in what they prove — and photograph-plus-self-report scores 0% because it produces no observation series to reconcile against. Radar therefore keeps the role VM0051 Appendix 4 assigns it: a per-plot cross-check lowering the frequency of field verification [4] (Supplementary S9.2).

10.2 Limitations

Seven limitations bound the claims, all stated in full in Supplementary S21; the two that condition the headline results are these. *Simulated farmers*: forgeries are injected into simulated trajectories, so a field pilot is still needed to measure behavioural response — above all whether farmers shift their *actual* pumping once the anchor is known, an attack that changes M rather than L and that our model cannot represent. *The attack is demonstrated in the water-scarce year only*: in 2025 most logs fall below the three-command minimum at which a timing attack is constructible — of 37 genuinely altered logs under the full Θ , all 37 come from 2024. That escarpment is itself a finding (MRV fraud risk is concentrated where water is scarce), but the attack results are not yet evidence about a wet season; under the reduced Θ twelve altered logs do come from 2025 and behave the same way — suggestive rather than conclusive. The remaining five are enumerated in S21.

11 Conclusion

We set out to make rice MRV nearly free by auditing a farmer-written irrigation log against soil-water physics. The first version of that layer reported no false accusation and complete detection; both numbers were artefacts of a field generated by the auditing model itself, and rebuilding it against an adversary that optimises the credited quantity changes the conclusion.

The point-model detector falsely accuses 26.7% of honest farmers; a proportional threshold and an intersection over a 150-element uncertainty set reduce that to 5.0% and 0/60 at a measured cost in detection (9/9 to 4/9 on omission fraud). Meanwhile a volume-preserving permutation of claimed times inflates creditable dry phases by 33% (paired, $p < 0.001$; 38% under the full uncertainty set, but on three logs only and thus not significant) while escaping the robust tier in 8 of 9 and 3 of 3 altered logs, and no physical configuration escapes that trade-off: delaying a claim both lengthens the dry phase and increases the deficit available to it, so profit and undetectability share a gradient.

The repair is not hydrological. Matching the claimed log against the gateway's signed, time-stamped command record catches 116/116 genuinely forged logs and accuses none of 60 honest ones, at no cost beyond a controller the companion work already deploys. What survives of the original premise: physics *constrains* a log, but only a signed device record *attests* it — near-zero-cost MRV is achievable only where a device stands between the farmer and the pump. Three previously published numbers are corrected accordingly ([49.0, 86.0] mm rather than 61 mm; an inert third filter; an inadmissible absolute gap test). Code, both uncertainty sets, all 281 audited logs and 46 unit tests are public.

Data, Code and Declarations

Code, both uncertainty sets, both adversaries and all 281 audited logs are at <https://github.com/mxuanvan02/TwinGate> (MIT licence); weather forcing is public ERA5 reanalysis via Open-Meteo. Funded by Hue University, Grant No. DHH2025-19-07. No competing interests. Ethics approval is not applicable (simulation and public data only).

Tóm tắt và từ khóa (Tiếng Việt)

Đo đạc, báo cáo, thẩm định (MRV) là nút thắt khiến mức giảm phát thải của lúa phát thải thấp chưa thành tín chỉ các-bon; giám sát chiếm ba phần tư chi phí. Vì mọi hành động quản lý nước đủ điều kiện cấp tín chỉ trong tưới khô ướt luân phiên (AWD) vốn là lệnh điều khiển, nhật ký nông dân tự ghi trông như bằng chứng gần như miễn phí, xếp hạng cao nhất theo Quyết định 4801/QĐ-BNNMT. Chúng tôi dựng lớp đồ trên mô hình sinh đôi số FAO-56 rồi tấn công nó; phần tấn công là đóng góp. Phát lại nhật ký với ERA5 công khai phơi bày độ sâu không đất nào chứa được, nhưng đối thủ giữ nguyên tổng lượng và từng độ sâu, chỉ dịch thời điểm lệnh, làm số pha khô được cấp tín chỉ tăng 33% (ghép cặp). Không lớp vật lý nào thoát hữu ích: bộ phát hiện mô hình điểm kết tội oan 26,7% nông dân trung thực, còn bộ bền vững giảm oan xuống 5,0% lại bỏ

lọt phần lớn ca dịch thời điểm. Trên 281 nhật ký tại ba trạm Đồng bằng sông Cửu Long, hai năm khí hậu, đối chiếu với bản ghi lệnh có chữ ký, dấu thời gian của gateway bắt 116/116 nhật ký giả, không oan ai. Vật lý không thể đứng một mình; dấu thời gian phải do thiết bị cấp, không phải nông dân.

Từ khóa: tưới khô ướt luân phiên (AWD), đo đạc báo cáo thẩm định (MRV), tín chỉ các-bon, lúa phát thải thấp, mô hình sinh đôi số.

Thông tin tác giả. Mai Xuân Văn¹, Nguyễn Tương Tri^{2,3,*}

¹ Trường Trung học phổ thông (THPT) Thuận Hóa, Trường Đại học Sư phạm, Đại học Huế; ² Viện Đào tạo Mở và Công nghệ thông tin, Đại học Huế; ³ Trường Đại học Sư phạm, Đại học Huế, 34 Lê Lợi, Huế, Việt Nam

Học hàm, học vị: (diễn); điện thoại: (diễn); E-mail: ntuongtri@hueuni.edu.vn; * Tác giả liên hệ: Nguyễn Tương Tri.

Nghiên cứu này được tài trợ bởi Đại học Huế trong đề tài mã số DHH2025-19-07.

Chữ ký tác giả:

References

- [1] Balaine N, Carrijo DR, Adviento-Borbe MAA, Linquist BA. Greenhouse Gases from Irrigated Rice Systems under Varying Severity of Alternate-Wetting and Drying Irrigation. *Soil Science Society of America Journal*. 2019;83(5):1533-41.
- [2] Yan X, Yagi K, Akiyama H, Akimoto H. Statistical analysis of the major variables controlling methane emission from rice fields. *Global Change Biology*. 2005;11(7):1131-41.
- [3] Bộ Nông nghiệp và Môi trường. Quyết định 4801/QĐ-BNNMT: Quy trình thí điểm Đo đạc, báo cáo, thẩm định (MRV) trong canh tác lúa chất lượng cao, phát thải thấp vùng Đồng bằng sông Cửu Long; 2025. Hanoi, Viet Nam, 14 November 2025.
- [4] Verra. VM0051 Improved Management in Rice Production Systems, Version 1.1; 2026. Effective 14 July 2026. <https://verra.org>.
- [5] Vanderheyden L, Gayot N, Van den Broeck G. What monitoring, reporting & verification (MRV) systems can reduce costs and enhance scalability of carbon farming? *Journal of Environmental Management*. 2026;401:128885.
- [6] Mai XV, Nguyen TT. A Digital-Twin Safety Gate for Irrigation Command Dispatch over Lossy Networks and Its Relevance to Emission MRV; 2026. Manuscript under review.
- [7] Allen RG, Pereira LS, Raes D, Smith M. Crop Evapotranspiration — Guidelines for Computing Crop Water Requirements. FAO Irrigation and Drainage Paper 56. Rome: Food and Agriculture Organization of the United Nations; 1998.
- [8] Hashemi MGZ, Abhishek A, Jalilvand E, Jayasinghe S, Andreadis KM, Siqueira P, et al. Assessing the impact of Sentinel-1 derived planting dates on rice crop yield modeling. *International Journal of Applied Earth Observation and Geoinformation*. 2022;114:103047.

- [9] Clauss K, Ottinger M, Kuenzer C. Mapping rice areas with Sentinel-1 time series and super-pixel segmentation. *International Journal of Remote Sensing*. 2017;39(5):1399-420.
- [10] Dela Cruz KMS, Ella VB, Suministrado DC, et al. A Low-Cost Wireless Sensor for Real-Time Monitoring of Water Level in Lowland Rice Field under Alternate Wetting and Drying. *Water*. 2022;14(24):4128.
- [11] Lovell RJ. Identifying Alternative Wetting and Drying (AWD) Adoption in the Vietnamese Mekong River Delta. *ISPRS International Journal of Geo-Information*. 2019;8(7):312.
- [12] Islam MR, Oyoshi K, Takeuchi W. Alternate Wetting and Drying Irrigated Rice Paddy Field Water Status Monitoring with ALOS-2 Three-Component Scattering Power Decomposition. *Remote Sensing*. 2026;18(8):1183.
- [13] Shah A, Singh R, Zaytar A, Tadesse GA, Robinson C, Tafti N, et al. Machine Learning for Sustainable Rice Production: Region-Scale Monitoring of Water-Saving Practices in Punjab, India. *arXiv preprint arXiv:250708605*. 2025.
- [14] Thủ tướng Chính phủ. Quyết định 1490/QĐ-TTg: Phê duyệt Đề án “Phát triển bền vững một triệu héc-ta chuyên canh lúa chất lượng cao và phát thải thấp gắn với tăng trưởng xanh vùng đồng bằng sông Cửu Long đến năm 2030”; 2023. Hanoi, Viet Nam, 27 November 2023.
- [15] Ministry of Agriculture and Environment of Viet Nam. More than 350,000 hectares of low-emission rice cultivated in Viet Nam; 2026. Programme progress, 11 MRV pilots in 2025, emission reduction of 3–4 t CO₂e per hectare, farmer income up 13.4%. Accessed 05/10/2026. <https://en.mae.gov.vn/more-than-350000-hectares-of-lowemission-rice-cultivated-in-viet-nam-9192.htm>.
- [16] Gonzalez T, Concepcion G. Why MRV is the missing link for low-emission rice in Southeast Asia; 2026. *Rice Today*, IRRI, 16 April 2026. ricetoday.irri.org.
- [17] Carrijo DR, Lundy ME, Linquist BA. Rice yields and water use under alternate wetting and drying irrigation: A meta-analysis. *Field Crops Research*. 2017;203:173-80.
- [18] Lampayan RM, Reyes RM, Singleton GR, Bouman BAM. Adoption and economics of alternate wetting and drying water management for irrigated lowland rice. *Field Crops Research*. 2015;170:95-108.
- [19] Enriquez Y, Yadav S, Evangelista G, et al. Disentangling Challenges to Scaling Alternate Wetting and Drying Technology for Rice Cultivation. *Frontiers in Sustainable Food Systems*. 2021;5.